

# LASSO

Lingyu Zhao

```
# Clear environment  
rm(list = ls())
```

# Loading Data and Exploratory Data Analysis

```
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

# Read in dataset
data = read.csv("https://raw.githubusercontent.com/Panta-Rhei-LZ/MDA_9159_Team_Bits_Project/main/Team_Bits_9159.csv")

# Remove price=0 entries
data = data[data$PRICE != 0, ]

# Remove rows with NA in all columns except 'YR_RMDL'
data = data %>% dplyr::filter(!if_any(-YR_RMDL, is.na))

# Remove not useful columns
data = data %>% dplyr::select(-SSL, -OBJECTID, -GIS_LAST_MOD_DTTM,
                             -QUALIFIED, -SALE_NUM, -BLDG_NUM,
                             -STYLE_D, -STRUCT_D, -GRADE_D,
                             -CNDTN_D, -EXTWALL_D, -ROOF_D,
                             -INTWALL_D, -USECODE, -HEAT_D,
                             -NUM_UNITS, -STRUCT)

head(data)
```

	BATHRM	HF_BATHRM	HEAT	AC	ROOMS	BEDRM	AYB	YR_RMDL	EYB	STORIES	
## 1	4		1	8	Y	12	6	1911	2021	1989	3.75
## 2	3		1	1	Y	13	5	1912	2009	1978	3.00
## 3	3		1	7	Y	6	4	1910	2022	1993	3.00
## 4	3		1	7	Y	11	4	1912	2000	1978	3.00
## 5	4		1	1	Y	11	5	1912	2007	1993	3.00
## 6	7		1	8	Y	16	7	1895	2014	1993	3.00
	SALEDATE	PRICE	GBA	STYLE	GRADE	CNDTN	EXTWALL	ROOF	INTWALL		
## 1	2019/08/19 04:00:00+00	3275000	6765	10	8	4	20	11	6		
## 2	1999/08/04 04:00:00+00	550000	2282	7	6	4	14	2	6		
## 3	2019/07/22 04:00:00+00	1700000	2016	7	6	4	14	6	6		
## 4	2021/10/27 04:00:00+00	1500000	2034	7	6	4	14	6	6		
## 5	2023/04/18 04:00:00+00	2232500	2655	7	6	5	14	2	6		
## 6	2013/12/30 05:00:00+00	1320000	2894	7	6	5	14	6	6		
	KITCHENS	FIREPLACES	LANDAREA								
## 1	1		6	2104							
## 2	2		3	936							
## 3	2		2	936							
## 4	2		2	988							
## 5	3		4	1674							
## 6	4		1	1674							

## Variable Explanation

We are dealing with housing data in this report, let me go over through the meanings behind each predictor:

1. PRICE: response
2. BATHRM: # bathrooms
3. HF\_BATHRM: # half bathrooms
4. HEAT: heating
5. AC: air conditioning
6. ROOMS: # rooms
7. BEDRM: # bedrooms
8. AYB: The earliest time the main portion of the building was built
9. YR\_RMDL: Year structure was remodelled
10. EYB: The year an improvement was built
11. STORIES: # stories in primary dwelling
12. SALEDATE: Date of sale
13. GBA: Gross building area in square feet
14. STYLE: House style
15. GRADE: House grade
16. CNDTN: House condition
17. EXTWALL: Exterior wall tyle
18. ROOF: Roof type
19. INTWALL: Interior wall type
20. KITCHENS: # kitchens
21. FIREPLACES: # fireplaces
22. LANDAREA: Land area of property in square feet

## NA Data

Now let us explore the percentage of missing data for each predictor:

```
missing_data = round(sapply(data, function(x) mean(is.na(x) * 100)), 3)
```

```
missing_data
```

##	BATHRM	HF_BATHRM	HEAT	AC	ROOMS	BEDRM	AYB
##	0.000	0.000	0.000	0.000	0.000	0.000	0.000
##	YR_RMDL	EYB	STORIES	SALEDATE	PRICE	GBA	STYLE
##	36.432	0.000	0.000	0.000	0.000	0.000	0.000
##	GRADE	CNDTN	EXTWALL	ROOF	INTWALL	KITCHENS	FIREPLACES
##	0.000	0.000	0.000	0.000	0.000	0.000	0.000
##	LANDAREA						
##	0.000						

From the R output above, observe that “YR\_RMDL: Year structure was remodeled” has around 36% missing data. A possible explanation for this could be: not all buildings were remodeled.

## Preprocessing

- Converted some predictors to numerical values:
  - AC: “Y” and “N” corresponds to “1” and “0”.
  - SALEDATE: Transform calendar format values in SALEDATE to numerical values using `as.Date()`.
- Created dummy variables for categorical predictors:

- These categorical variables include: “AC”, “HEAT”, “STYLE”, “GRADE”, “CNDTN”, “EXTWALL”, “ROOF” and “INTWALL”.
- Introduced a few new variables:
  - SALE\_YEAR: The year that the house was sold, it is derived from SALEDATE.
  - SALE\_AYB\_DIFF: The difference between the year sold and the year built.
  - SALE\_EYB\_DIFF: The difference between the year sold and the year an improvement was applied.
  - SALE\_RMDL\_DIFF: The difference between the year sold and the year structure was remodeled.

```
library(lubridate)

##
## Attaching package: 'lubridate'
## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union

library(fastDummies)

# Transform Yes/No for having AC to numerical values
data$AC = ifelse(data$AC == 'Y', 1, 0)

# Add SALEYEAR
data$SALE_YEAR = year(ymd_hms(data$SALEDATE))

# Add SALEYEAR and AYB diff
data$SALE_AYB_DIFF = data$SALE_YEAR - data$AYB

# Add SALEYEAR and EYB diff
data$SALE_EYB_DIFF = data$SALE_YEAR - data$EYB

# Add SALEYEAR and YR_RMDL diff
data$SALE_RMDL_DIFF = data$SALE_YEAR - data$YR_RMDL

# Convert SALEDATE column to numeric values
data$SALEDATE = as.numeric(as.Date(data$SALEDATE))

# Replace NA with column median
data = data.frame(lapply(data, function(column) {
  column_median = median(column, na.rm = TRUE)
  column[is.na(column)] = column_median
  column
}))

set.seed(9159)

# Temporarily storing current data structure for plotting in later steps
temp_data = data[sample(nrow(data), 600),]

# Create dummy variables for categorical predictors
data = dummy_cols(
  data,
  select_columns = c("HEAT", "STYLE", "GRADE", "CNDTN",
    "EXTWALL", "ROOF", "INTWALL", "AC"),
```

```

remove_selected_columns = TRUE,
remove_first_dummy = TRUE
)

# Define box-cox and inverse box-cox transformation
powerfun = function(y, lambda) {
  if (lambda == 0) {
    return(log(y))
  } else {
    return((y^lambda - 1) / lambda)
  }
}

inv_powerfun = function(y_transformed, lambda) {
  if (lambda == 0) {
    return(exp(y_transformed))
  } else {
    return((lambda * y_transformed + 1)^(1/lambda))
  }
}

```

## Data for Training and Validating

```

set.seed(9159)

# Randomly sample 600 data entries for our project
clean_data = data[sample(nrow(data), 600),]

data_train = clean_data[1:500, ] # First 500 rows for training
data_valid = clean_data[501:600, ] # Last 100 rows for validation

```

## Response Variable Transformation

```

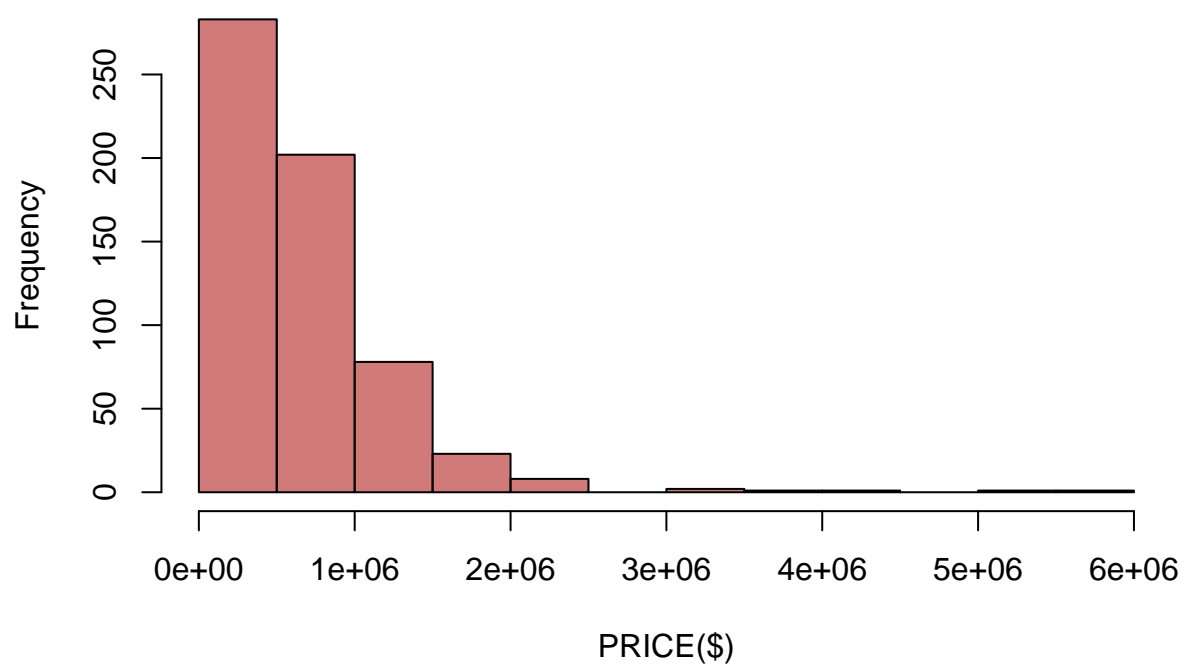
library(lmtest)

## Loading required package: zoo
##
## Attaching package: 'zoo'
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric

# Plot histogram for untransformed response variable `PRICE`
hist(clean_data$PRICE, col=adjustcolor('firebrick',alpha=0.6),
      xlab='PRICE($)', ylab='Frequency',
      main='Histogram of Untransformed PRICE')

```

## Histogram of Untransformed PRICE




```
shapiro.test(clean_data$PRICE)
```

```
##  
##  Shapiro-Wilk normality test  
##  
## data:  clean_data$PRICE  
## W = 0.76073, p-value < 2.2e-16
```

```
library(MASS)
```

```
##  
## Attaching package: 'MASS'  
## The following object is masked from 'package:dplyr':  
##  
##      select
```


```
boxcox(lm(PRICE~., data=clean_data),lambda=seq(0,0.5,by=0.05))
```



```
lambda = 0.35
```

```
# Plot histogram for log-transformed response variable `PRICE`
hist(powerfun(clean_data$PRICE, lambda), col=adjustcolor('firebrick',alpha=0.6),
     xlab='PRICE($)', ylab='Frequency',
     main='Histogram of Log-Transformed PRICE')
```

## Histogram of Log-Transformed PRICE



```
shapiro.test(powerfun(clean_data$PRICE, lambda))
```

```
##  
##  Shapiro-Wilk normality test  
##  
## data:  powerfun(clean_data$PRICE, lambda)  
## W = 0.98117, p-value = 5.504e-07
```

```
trans_PRICE = powerfun(clean_data$PRICE, lambda)
```


## Data Visualization

```
library(corrplot)
```

```
## corrplot 0.95 loaded
```

```
cor_temp_data = cor(temp_data, use='complete.obs', method='pearson')  
corrplot(cor_temp_data, method='color', type='lower', tl.col='black', tl.cex=0.5)
```





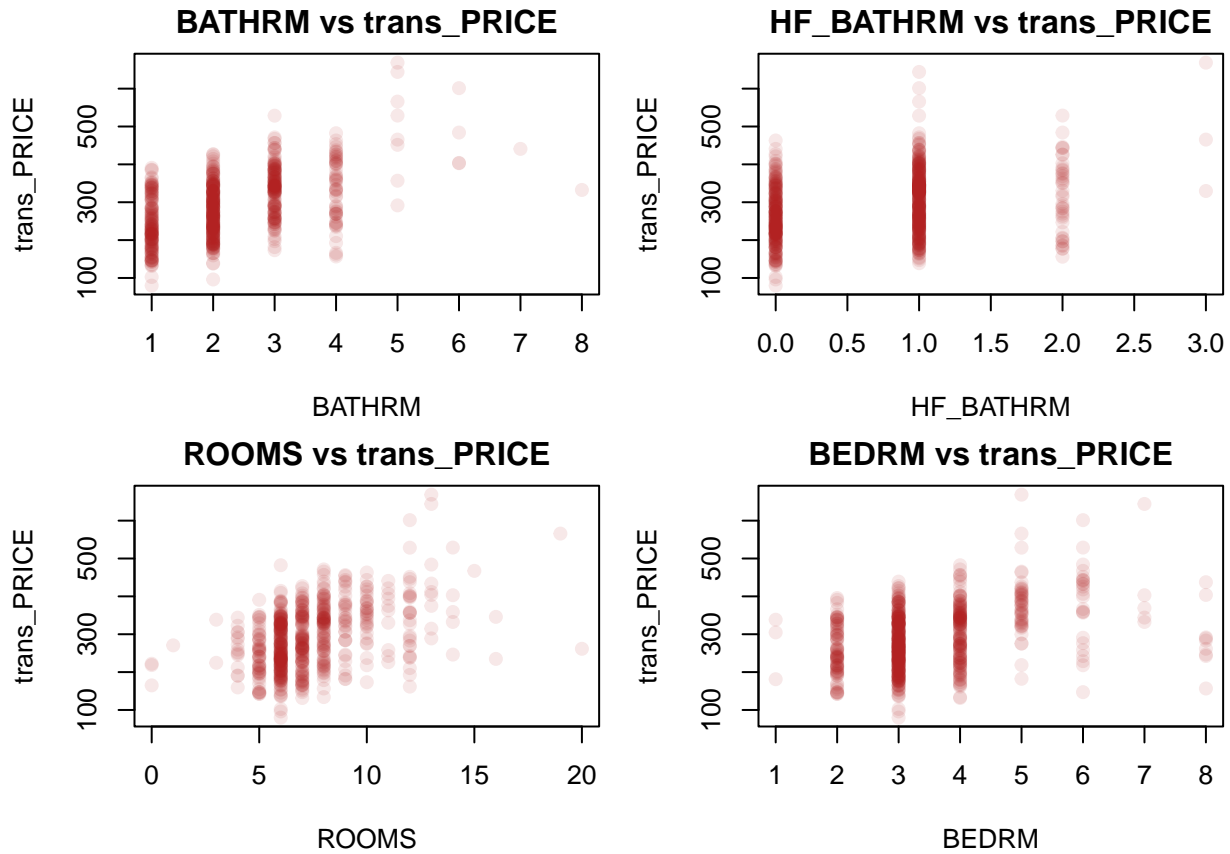
```
par(mfrow = c(2, 2), mar = c(4, 4, 2, 1))

plot(clean_data$BATHRM, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "BATHRM", ylab = "trans_PRICE",
     main = "BATHRM vs trans_PRICE")

plot(clean_data$HF_BATHRM, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "HF_BATHRM", ylab = "trans_PRICE",
     main = "HF_BATHRM vs trans_PRICE")

plot(clean_data$ROOMS, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "ROOMS", ylab = "trans_PRICE",
     main = "ROOMS vs trans_PRICE")

plot(clean_data$BEDRM, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "BEDRM", ylab = "trans_PRICE",
     main = "BEDRM vs trans_PRICE")
```



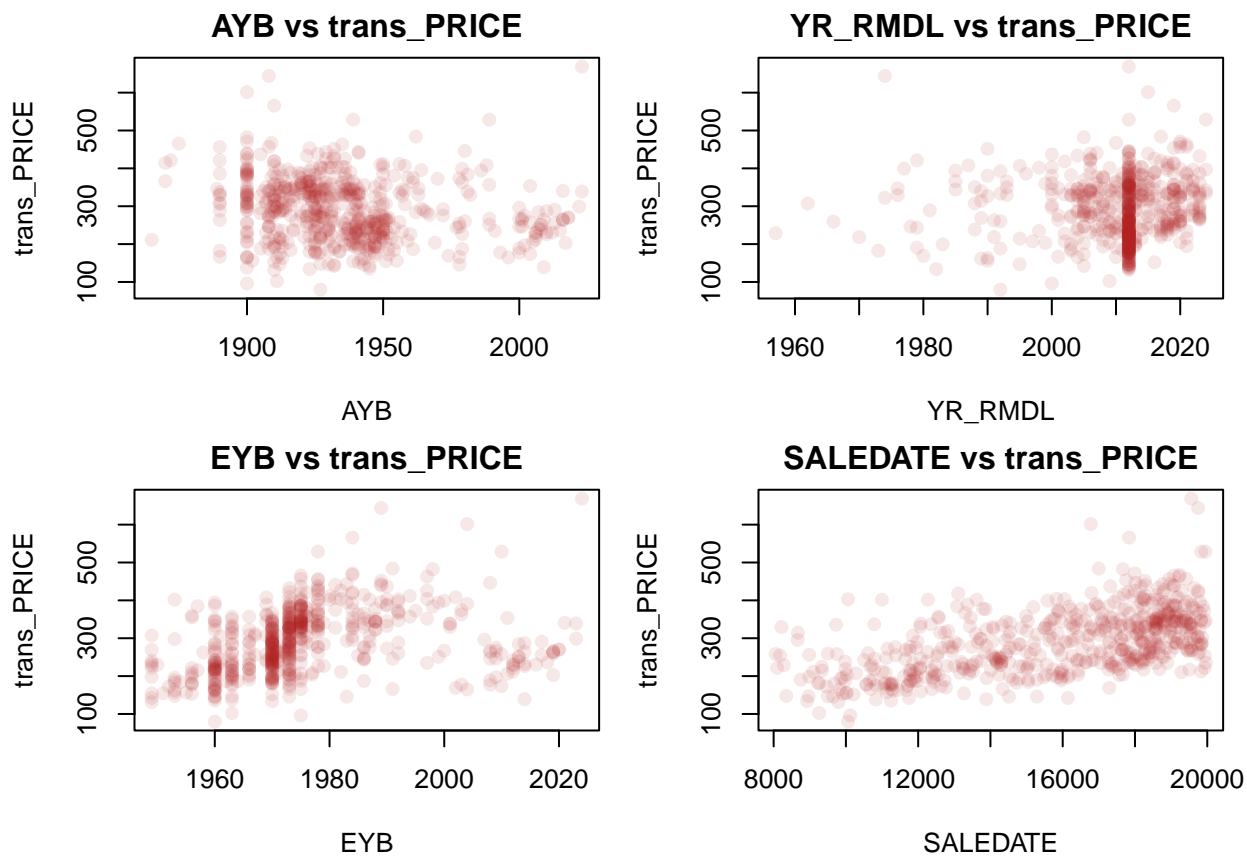
```
par(mfrow = c(2, 2), mar = c(4, 4, 2, 1))

plot(clean_data$AYB, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "AYB", ylab = "trans_PRICE",
     main = "AYB vs trans_PRICE")

plot(clean_data$YR_RMDL, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "YR_RMDL", ylab = "trans_PRICE",
     main = "YR_RMDL vs trans_PRICE")

plot(clean_data$EYB, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "EYB", ylab = "trans_PRICE",
     main = "EYB vs trans_PRICE")

plot(clean_data$SALEDATE, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "SALEDATE", ylab = "trans_PRICE",
     main = "SALEDATE vs trans_PRICE")
```



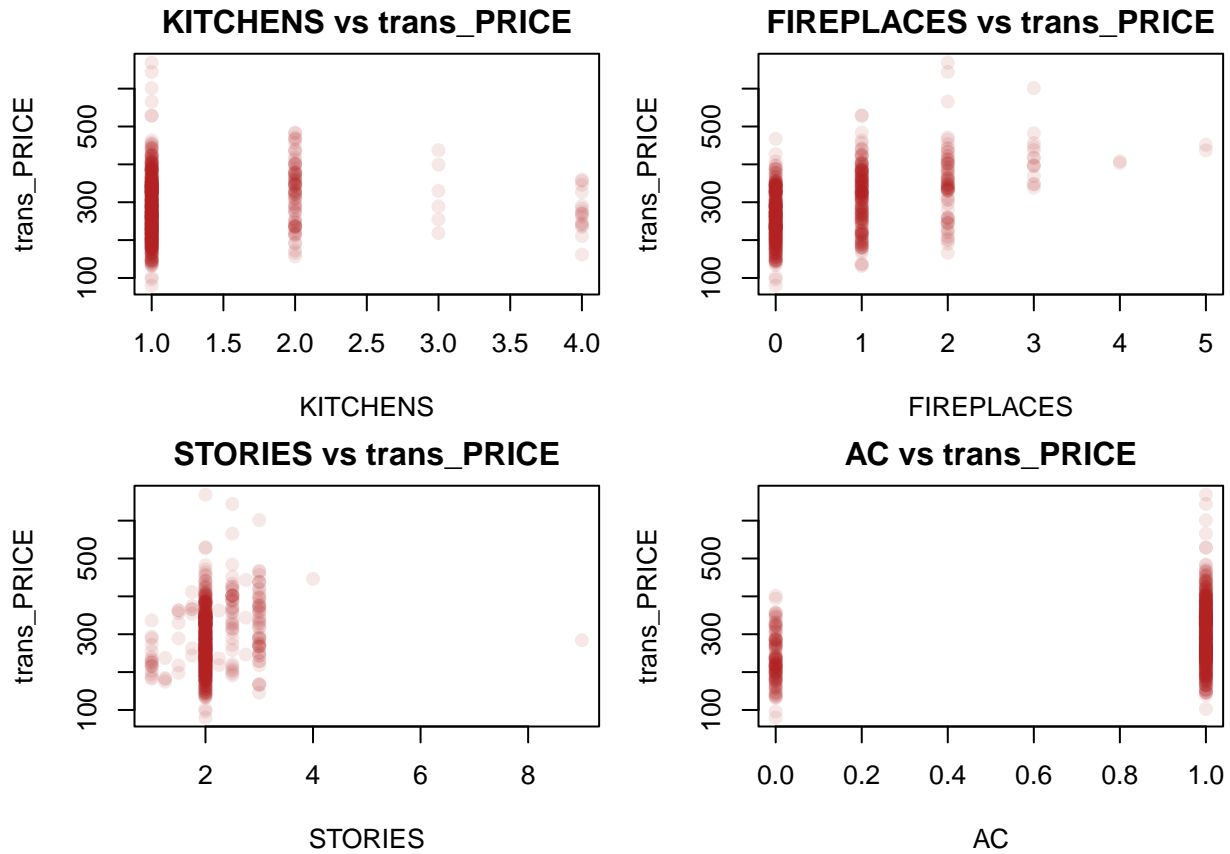
```
par(mfrow = c(2, 2), mar = c(4, 4, 2, 1))

plot(clean_data$KITCHENS, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "KITCHENS", ylab = "trans_PRICE",
     main = "KITCHENS vs trans_PRICE")

plot(clean_data$FIREPLACES, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "FIREPLACES", ylab = "trans_PRICE",
     main = "FIREPLACES vs trans_PRICE")

plot(clean_data$STORIES, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "STORIES", ylab = "trans_PRICE",
     main = "STORIES vs trans_PRICE")

plot(clean_data$AC, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "AC", ylab = "trans_PRICE",
     main = "AC vs trans_PRICE")
```



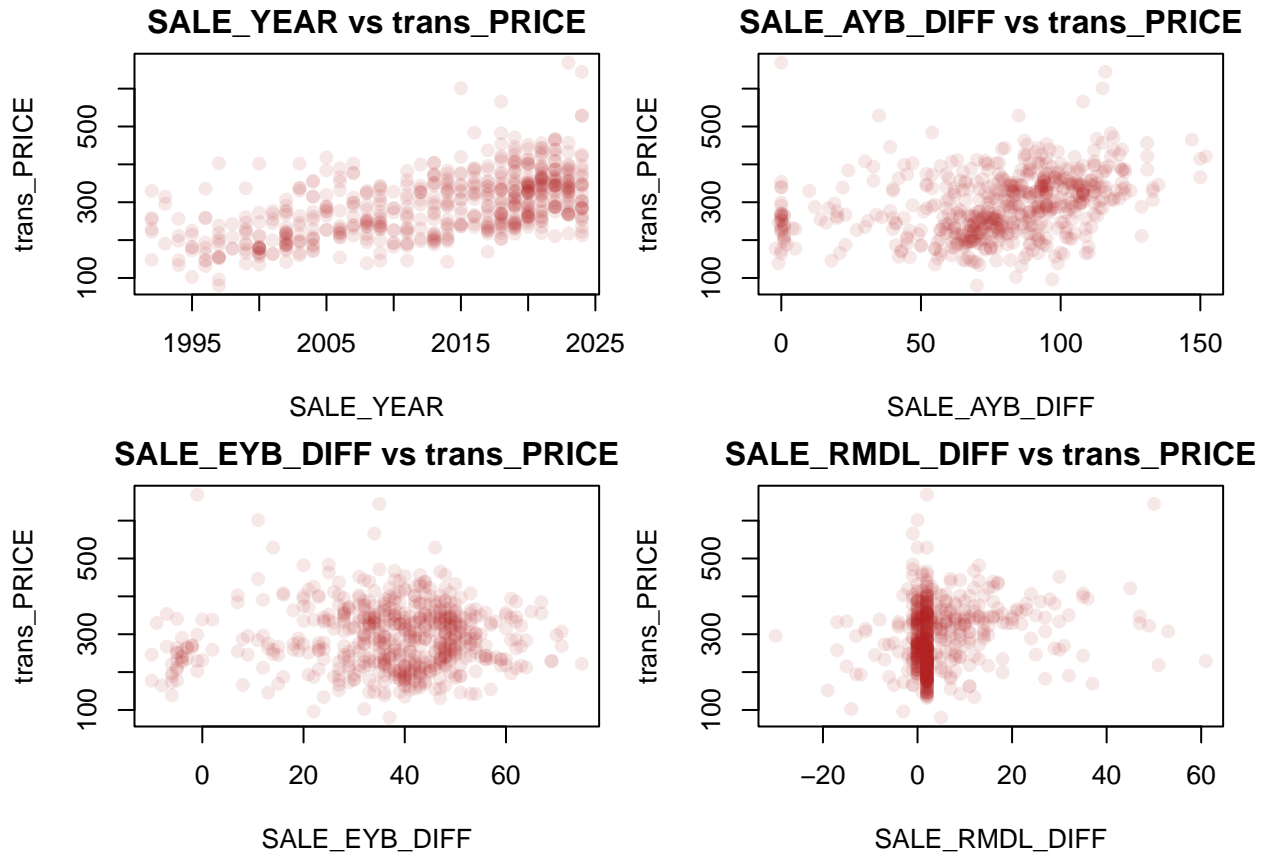
```
par(mfrow = c(2, 2), mar = c(4, 4, 2, 1))

plot(clean_data$SALE_YEAR, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "SALE_YEAR", ylab = "trans_PRICE",
     main = "SALE_YEAR vs trans_PRICE")

plot(clean_data$SALE_AYB_DIFF, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "SALE_AYB_DIFF", ylab = "trans_PRICE",
     main = "SALE_AYB_DIFF vs trans_PRICE")

plot(clean_data$SALE_EYB_DIFF, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "SALE_EYB_DIFF", ylab = "trans_PRICE",
     main = "SALE_EYB_DIFF vs trans_PRICE")


plot(clean_data$SALE_RMDL_DIFF, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "SALE_RMDL_DIFF", ylab = "trans_PRICE",
     main = "SALE_RMDL_DIFF vs trans_PRICE")
```



```
par(mfrow = c(1, 2), mar = c(4, 4, 2, 1))

plot(clean_data$GBA, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "GBA", ylab = "trans_PRICE",
     main = "GBA vs trans_PRICE")

plot(clean_data$LANDAREA, trans_PRICE,
     pch = 19, col = adjustcolor('firebrick', alpha = 0.1),
     xlab = "LANDAREA", ylab = "trans_PRICE",
     main = "LANDAREA vs trans_PRICE")
```



## Model Building and Analysis

### Model Training without Non-Linear Predictors

```
library(glmnet)

## Loading required package: Matrix
## Loaded glmnet 4.1-8
X_train = model.matrix(powerfun(PRICE, lambda)~., data=data_train)[,-1]
X_test = model.matrix(powerfun(PRICE, lambda)~., data=data_valid)[,-1]
y_train = powerfun(data_train$PRICE, lambda)
y_test = powerfun(data_valid$PRICE, lambda)

cv_lasso = cv.glmnet(X_train, y_train, alpha=1)
best_lasso_lambda = cv_lasso$lambda.min

lasso_model = glmnet(X_train, y_train, alpha=1, lambda=best_lasso_lambda)

# Check non-zero LASSO coefficients
predict(lasso_model, type="coefficients", s=best_lasso_lambda)

## 116 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept)  -3.262452e+02
## BATHRM       4.588097e+00
```

## HF_BATHRM	3.904705e+00
## ROOMS	1.149480e+00
## BEDRM	5.370663e-01
## AYB	.
## YR_RMDL	.
## EYB	1.674562e-01
## STORIES	.
## SALEDATE	1.190030e-02
## GBA	1.663851e-02
## KITCHENS	.
## FIREPLACES	1.590030e+01
## LANDAREA	1.705915e-03
## SALE_YEAR	.
## SALE_AYB_DIFF	4.573229e-01
## SALE_EYB_DIFF	.
## SALE_RMDL_DIFF	.
## HEAT_1	-8.221947e+00
## HEAT_2	-1.746379e+01
## HEAT_3	.
## HEAT_4	.
## HEAT_5	.
## HEAT_6	.
## HEAT_7	.
## HEAT_8	2.167639e+01
## HEAT_9	.
## HEAT_10	.
## HEAT_11	.
## HEAT_12	.
## HEAT_13	3.450422e-01
## STYLE_1	-6.345856e+00
## STYLE_2	.
## STYLE_3	.
## STYLE_4	.
## STYLE_5	.
## STYLE_6	.
## STYLE_7	.
## STYLE_8	.
## STYLE_9	.
## STYLE_10	.
## STYLE_11	.
## STYLE_12	.
## STYLE_13	.
## STYLE_14	.
## STYLE_15	.
## STYLE_94	.
## STYLE_99	.
## GRADE_1	.
## GRADE_2	.
## GRADE_3	-3.041416e+01
## GRADE_4	-9.805268e+00
## GRADE_5	.
## GRADE_6	1.983899e+01
## GRADE_7	2.371386e+01
## GRADE_8	2.054186e+01

## GRADE_9	7.057614e+01
## GRADE_10	.
## GRADE_11	.
## GRADE_12	.
## CNDTN_1	.
## CNDTN_2	-1.550982e+01
## CNDTN_3	-1.415970e+01
## CNDTN_4	.
## CNDTN_5	1.617908e+01
## CNDTN_6	.
## EXTWALL_1	.
## EXTWALL_2	.
## EXTWALL_3	.
## EXTWALL_4	.
## EXTWALL_5	3.900062e+00
## EXTWALL_6	.
## EXTWALL_7	.
## EXTWALL_8	.
## EXTWALL_10	1.681968e+01
## EXTWALL_11	.
## EXTWALL_12	.
## EXTWALL_13	.
## EXTWALL_14	.
## EXTWALL_15	.
## EXTWALL_16	.
## EXTWALL_17	.
## EXTWALL_18	.
## EXTWALL_19	-1.514672e+01
## EXTWALL_20	-1.910498e+01
## EXTWALL_21	.
## EXTWALL_22	-5.043914e-01
## EXTWALL_23	.
## EXTWALL_24	.
## ROOF_1	-2.751832e+00
## ROOF_2	.
## ROOF_3	.
## ROOF_4	-5.339574e+00
## ROOF_5	.
## ROOF_6	8.714045e+00
## ROOF_7	.
## ROOF_8	.
## ROOF_9	.
## ROOF_10	.
## ROOF_11	.
## ROOF_12	.
## ROOF_13	.
## ROOF_14	.
## ROOF_15	.
## INTWALL_1	.
## INTWALL_2	-1.630494e+01
## INTWALL_3	.
## INTWALL_4	.
## INTWALL_5	.
## INTWALL_6	.



```
## INTWALL_7      .
## INTWALL_8      1.874432e+01
## INTWALL_9      .
## INTWALL_10     .
## INTWALL_11     -1.242621e+00
## AC_1           9.564972e+00

# Predicted values on training data
pred_train = predict(lasso_model, newx=X_train, s=best_lasso_lambda)

# Compute training SSE and SST
SSE = sum((y_train - pred_train)^2)
SST = sum((y_train - mean(y_train))^2)

# R-squared
R2 = 1 - SSE / SST

# Adjusted R-squared
n_train = length(y_train)
p_train = sum(coef(lasso_model, s=best_lasso_lambda) != 0)
Adjusted_R2 = 1 - (1 - R2) * (n_train) / (n_train - p_train)

cat("R-squared:", R2, '\n')
```

```
## R-squared: 0.8325244
```

```
cat("Adjusted R-squared:", Adjusted_R2, '\n')
```

```
## Adjusted R-squared: 0.8191409
```

## Compute Loss Metrics

```
# Predicted values are transformed by powerfun(PRICE, lambda)
valid_pred = predict(lasso_model, s=best_lasso_lambda, newx=X_test)

# Using inv_powerfun to convert back to original scale
inv_valid_pred = inv_powerfun(valid_pred, lambda)
```

## Compute MSE, RMSE and RMSLE

```
# Compute metrics on validation dataset
mse = mean((data_valid$PRICE - inv_valid_pred)^2)
rmse = sqrt(mse)
rmsle = sqrt(mean((log(data_valid$PRICE) - log(inv_valid_pred))^2))

cat("MSE:", mse, '\n')
```

```
## MSE: 145347824793
```

```
cat("RMSE:", rmse, '\n')
```

```
## RMSE: 381245.1
```

```
cat("RMSLE:", rmsle, '\n')
```

```
## RMSLE: 0.4067307
```

## Model Training including Non-Linear Predictors

```
X_train2 = model.matrix(
  powerfun(PRICE, lambda)~. + poly(BATHRM, 4) +
  poly(HF_BATHRM, 3) + poly(BEDRM, 3) + poly(EYB, 2),
  data=data_train
)[,-1]
X_test2 = model.matrix(
  powerfun(PRICE, lambda)~. + poly(BATHRM, 4) +
  poly(HF_BATHRM, 3) + poly(BEDRM, 3) + poly(EYB, 2),
  data=data_valid
)[,-1]

cv_lasso2 = cv.glmnet(X_train2, y_train, alpha=1)
best_lasso_lambda2 = cv_lasso2$lambda.min

lasso_model2 = glmnet(X_train2, y_train, alpha=1, lambda=best_lasso_lambda2)

# Check non-zero LASSO coefficients
predict(lasso_model2, type="coefficients", s=best_lasso_lambda2)

## 128 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept)    15.955213972
## BATHRM         1.381972647
## HF_BATHRM      2.346109349
## ROOMS          0.740200132
## BEDRM          0.174508299
## AYB            .
## YR_RMDL        .
## EYB            0.002131434
## STORIES        .
## SALEDATE       0.011860373
## GBA            0.019429632
## KITCHENS       .
## FIREPLACES     15.367339325
## LANDAREA       0.001565030
## SALE_YEAR      .
## SALE_AYB_DIFF  0.384952588
## SALE_EYB_DIFF  .
## SALE_RMDL_DIFF 0.074067132
## HEAT_1         -8.582266170
## HEAT_2        -20.013982805
## HEAT_3         .
## HEAT_4         .
## HEAT_5         .
## HEAT_6         .
## HEAT_7         .
## HEAT_8         20.589172254
## HEAT_9         .
## HEAT_10        .
## HEAT_11        .
## HEAT_12        .
## HEAT_13        .
## STYLE_1        -7.292723097
```

## STYLE_2	.
## STYLE_3	.
## STYLE_4	.
## STYLE_5	.
## STYLE_6	0.069959413
## STYLE_7	.
## STYLE_8	.
## STYLE_9	.
## STYLE_10	.
## STYLE_11	.
## STYLE_12	.
## STYLE_13	.
## STYLE_14	.
## STYLE_15	.
## STYLE_94	.
## STYLE_99	.
## GRADE_1	.
## GRADE_2	.
## GRADE_3	-28.247158831
## GRADE_4	-9.076734247
## GRADE_5	.
## GRADE_6	18.695135929
## GRADE_7	26.311045594
## GRADE_8	22.396547342
## GRADE_9	87.853112053
## GRADE_10	.
## GRADE_11	.
## GRADE_12	.
## CNDTN_1	.
## CNDTN_2	-11.263890469
## CNDTN_3	-13.067309907
## CNDTN_4	.
## CNDTN_5	16.768515233
## CNDTN_6	.
## EXTWALL_1	.
## EXTWALL_2	.
## EXTWALL_3	.
## EXTWALL_4	.
## EXTWALL_5	4.059318615
## EXTWALL_6	.
## EXTWALL_7	.
## EXTWALL_8	.
## EXTWALL_10	12.408006368
## EXTWALL_11	.
## EXTWALL_12	.
## EXTWALL_13	.
## EXTWALL_14	.
## EXTWALL_15	.
## EXTWALL_16	.
## EXTWALL_17	.
## EXTWALL_18	.
## EXTWALL_19	-14.083943081
## EXTWALL_20	-25.700326516
## EXTWALL_21	.

```

## EXTWALL_22      .
## EXTWALL_23      .
## EXTWALL_24      .
## ROOF_1          -3.091342746
## ROOF_2          .
## ROOF_3          .
## ROOF_4          -6.231959891
## ROOF_5          .
## ROOF_6          8.642179087
## ROOF_7          .
## ROOF_8          .
## ROOF_9          .
## ROOF_10         .
## ROOF_11         .
## ROOF_12         .
## ROOF_13         .
## ROOF_14         .
## ROOF_15         .
## INTWALL_1       .
## INTWALL_2       -14.672796892
## INTWALL_3       .
## INTWALL_4       .
## INTWALL_5       .
## INTWALL_6       .
## INTWALL_7       .
## INTWALL_8       22.005492132
## INTWALL_9       .
## INTWALL_10      .
## INTWALL_11      -0.762410072
## AC_1            6.946199933
## poly(BATHRM, 4)1  51.550307002
## poly(BATHRM, 4)2 -63.842247358
## poly(BATHRM, 4)3  .
## poly(BATHRM, 4)4  .
## poly(HF_BATHRM, 3)1 14.115285178
## poly(HF_BATHRM, 3)2  .
## poly(HF_BATHRM, 3)3  8.445482535
## poly(BEDRM, 3)1   13.909364743
## poly(BEDRM, 3)2   -52.025933361
## poly(BEDRM, 3)3   .
## poly(EYB, 2)1     36.128865397
## poly(EYB, 2)2     -91.500056357

# Predicted values on training data
pred_train = predict(lasso_model2, newx=X_train2, s=best_lasso_lambda2)

# Compute training SSE and SST
SSE = sum((y_train - pred_train)^2)
SST = sum((y_train - mean(y_train))^2)

# R-squared
R2 = 1 - SSE / SST

# Adjusted R-squared

```

```

n_train = length(y_train)
p_train = sum(coef(lasso_model2, s=best_lasso_lambda2) != 0)
Adjusted_R2 = 1 - (1 - R2) * (n_train) / (n_train - p_train)

cat("R-squared:", R2, '\n')

```

```
## R-squared: 0.8387716
```

```
cat("Adjusted R-squared:", Adjusted_R2, '\n')
```

```
## Adjusted R-squared: 0.8228259
```

### Compute Loss Metrics

```

# Predicted values are transformed by powerfun(PRICE, lambda)
valid_pred = predict(lasso_model2, s=best_lasso_lambda2, newx=X_test2)

# Using inv_powerfun to convert back to original scale
inv_valid_pred = inv_powerfun(valid_pred, lambda)

```

### Compute MSE, RMSE and RMSLE

```

# Compute metrics on validation dataset
mse = mean((data_valid$PRICE - inv_valid_pred)^2)
rmse = sqrt(mse)
rmsle = sqrt(mean((log(data_valid$PRICE) - log(inv_valid_pred))^2))

```

```
cat("MSE:", mse, '\n')
```

```
## MSE: 143750464990
```

```
cat("RMSE:", rmse, '\n')
```

```
## RMSE: 379144.4
```

```
cat("RMSLE:", rmsle, '\n')
```

```
## RMSLE: 0.4342245
```